

Contribution Title^{*}

First Author¹[0000–1111–2222–3333] and Fahiem Bacchus^{2,3}[1111–2222–3333–4444]

¹ Princeton University, Princeton NJ 08544, USA

² Springer Heidelberg, Tiergartenstr. 17, 69121 Heidelberg, Germany
lncs@springer.com

<http://www.springer.com/gp/computer-science/lncs>

³ ABC Institute, Rupert-Karls-University Heidelberg, Heidelberg, Germany
{abc,lncs}@uni-heidelberg.de

1 Introduction

2 Preliminaries

3 Find i-uip Clause

Classical CDCL SAT solvers analyze each conflict with 1-UIP resolution learning scheme to derive an asserting clause C_1 . Even though prior studies shows that further resolutions on C_1 against the assertion trail ϕ cannot reduce the clause's literal block distance (LBD), the resolutions could potentially reduce the clause's size. Clause size is an important quality measurement because a smaller clause 1) consumes less memory, 2) requires less steps to force a literal and 3) decreases the size of future conflict clauses.

Inspired by both LBD and clause size, i-UIP resolution learning attempts to resolve away literals in C_1 against the assertion trail ϕ to minimize the clause size without increasing the clause's LBD. The goal of i-UIP learning is to find a clause C_i whose literals are either the unique implication points in their respective decision levels or are directly implied by literals from a foreign decision level. Alg. 1 is a pseudo-code implementation of i-UIP learning.

The algorithm **i-UIP** computes C_i from (line 2 to 20), and then returns the smaller clause between C_1 and C_i (line 21 to 6). The algorithm first initializes C_i with the minimized C_1 . Next, the algorithm iterate through the decision levels of C_i in descending order (line 4) and tries to find the unique implication point in each level (line 6). Since **i-UIP** needs to preserve C_i 's LBD during resolutions, before resolving away a target literal p , the algorithm preemptively checks whether the reasoning clause of $\neg p$ contains any literal q from an unseen decision level (line 9). If such a literal q exists, then the algorithm cannot find the unique implication point (UIP) at the current level i because resolving away p will introduce q into the clause and increases LBD. One solution is to abandon level i by reverting back to the state before any literal at level i is resolved away (line 10), and then move on to the next level (line 11). We denote **i-UIP** learning with this solution as **Pure-i-UIP**. The final C_i produced by **Pure-i-UIP**

^{*} Supported by organization x.

Algorithm 1 i-UIP**Require:** C_1 is a valid and minimized 1-UIP clause**Require:** ϕ is a valid assertion trail

```

1: procedure i-UIP( $C_1, \phi$ )
2:    $C_i \leftarrow C_1$  ▷ initialize i-UIP clause
3:   DecisionLvs  $\leftarrow$  Decision levels in  $C_1$  in descending order
4:   for  $i \in$  DecisionLvs do
5:      $L_i \leftarrow \{l \mid \text{level}(l) = i \wedge l \in C_i\}$ 
6:     while  $|\text{unmarked}(L_i)| > 1$  do
7:        $p \leftarrow$  lit with the highest trail position in  $L_i$ 
8:       if  $\exists q \in \text{Reason}(\neg p, \phi) \cdot \text{level}(q) \notin$  DecisionLvs then
9:         if Pure-i-UIP then
10:           Unresolve literals at level  $i$ 
11:           Go to the next decision level  $i$ 
12:         else if Min-i-UIP then
13:           Mark( $p$ )
14:         end if
15:       else
16:          $C_i \leftarrow (C_1) \bowtie \text{Reason}(\neg p, \phi)$ 
17:         Update( $L_i$ )
18:       end if
19:     end while
20:   end for
21:   if  $|C_i| < |C_1|$  then
22:     return  $C_i$ 
23:   else
24:     return  $C_1$ 
25:   end if
26: end procedure

```

will contains exactly one literal for levels where the LBD-persevering UIP exist. However, **Pure-i-UIP** does not minimize the number of literals in the decision levels where the LBD-preserving UIP does not exist.

When **i-UIP** cannot resolve away a literal p without increasing C_i 's LBD, instead of skipping the entire decision level i , a more practical solution **Min-i-UIP** will mark and keep p in C_i (line 13), and continue resolutions at level i . **Min-i-UIP** will ignore all the unsolvable literals and find the "local" unique implication point at every decision level. The algorithm terminates when there is exactly one unmarked literal left for each decision level of C_i , representing the local unique implication points.

The core algorithm of **i-UIP** as a clause reduction technique is simple. However, to make the algorithm a practical learning scheme that is compatible with modern SAT solvers, a number of issues need to be addressed. The rest of the section details the key optimizations and augmentations of **i-UIP** as a practical clause learning scheme.

Algorithm 2 Control-i-UIP**Require:** C_1 is a valid 1-UIP clause**Require:** $t_{gap} \geq 0$ is a dynamically calculated gap threshold

```

1: procedure CONTROL-i-UIP( $C_1, t_{gap}$ )
2:    $Gap \leftarrow |C_1| - LBD(C_1)$ 
3:   if  $Gap > t_{gap}$  then
4:      $C_i \leftarrow \mathbf{i-UIP}(C_1, \phi)$ 
5:      $\mathbf{I-UIP-Greedy}(C_i, C_1)$   $\triangleright$  additional clause selection policy, see sec 3.3
6:     if  $|C_i| < |C_1|$  then
7:        $Succeed \leftarrow Succeed + 1$ 
8:     end if
9:      $Attempted \leftarrow Attempted + 1$ 
10:  end if
11: end procedure

```

3.1 Control i-UIP Learning

This simple and greedy **i-UIP** learning scheme can produce significantly smaller clause. However, when **i-UIP** does not reduce the clause size, the cost of the additional resolution steps will hurt the solver's performance. Since resolution cannot reduce a clause's LBD, The maximum size reduction from **i-UIP** is the difference between the clause's size and LBD, denote as the clause's gap value ($Gap(C_1) = |C_1| - LBD(C_1)$). For an 1-UIP clause with a small Gap, applying **i-UIP** is unlikely to achieve cost effective results. Therefore, we propose a heuristic based approach to enable and disable **i-UIP** learning based on input clause's Gap.

Alg. 2 compares the input C_1 's Gap against a floating target threshold t_{gap} (line 3). The threshold t_{gap} represent the expected minimal Gap required for C_1 to achieve a predetermined success rate (80%) from performing **i-UIP** learning.

The gap threshold t_{gap} is initialized to 0, and is updated at every restart based on **i-UIP**'s success rate from the previous restart interval. More specifically, the algorithm collects the statistics of the number of **i-UIP** learning attempted (line 9 in alg. 2) and the number of attempts succeed (line 7) for each restart interval, and use them to calculate the success rate. If the success rate is below 80, the threshold t_{gap} is increased to restrict **i-UIP** for the next restart interval. On the other hand, the threshold is decreased to encourage more aggressive **i-UIP** learning for the next restart interval.

$$t_{gap} = \begin{cases} t_{gap} + 1 & \text{if } \frac{Succeed}{Attempted} < 0.8 \\ \max(t_{gap} - 1, 0), & \text{otherwise} \end{cases}$$

3.2 Early stop i-UIP Learning

At any point of **i-UIP** learning, if the number of marked literals in C_i (literals which are forced into the clause to preserve LBD) exceeds the input clause C_1 's

Algorithm 3 i-UIP-Greedy

Require: C_i is a valid i-UIP clause

Require: C_1 is a valid 1-UIP clause

```

1: procedure i-UIP-Greedy( $C_i, C_1$ )
2:   if  $|C_i| < |C_1| \wedge (\text{AvgVarAct}(C_i) > \text{AvgVarAct}(C_1))$  then
3:     return  $C_i$ 
4:   else
5:     return  $C_1$ 
6:   end if
7: end procedure

```

Gap, we can abort **i-UIP** learning for C_1 because the size of the final C_i is at least the size of C_1 . The early stopping rule prevents solver from wasting time on traversing a large implication graph when **i-UIP** has already failed.

3.3 Greedy Active Clause Selection

Even though **i-UIP** learning can reduce the size of the learnt clause C_i , but it may introduce literals with low variable activity into C_i . Inactive literals prevents the clause from being asserted to force literal implication. Therefore, a practical clause learning scheme should consider both size and variable activity. We propose an optional extension **i-UIP-Greedy** to filter out inactive C_1 .

After computing the C_i , **i-UIP-Greedy** compares both the size and the average variable activity for C_1 and C_i (alg. 3 at line 2). The algorithm learns C_i if the clause has smaller size and higher average variable activity.

3.4 Adjust Variable Activity

Two popular branching heuristics for modern SAT solvers are VSIDS and LBR. Both heuristics increase the variable activity for all literals involved in resolutions during 1-UIP learning. Since **i-UIP** extends 1-UIP with deeper resolutions against the trail, the variables activities for the fresh literals involved in the additional resolution steps need to be adjusted as well. We propose two optional schemes for adjusting variable activities, **i-UIP-Inclusive** and **i-UIP-Exclusive**.

After learning C_1 from 1-UIP scheme, **i-UIP-Inclusive** collects all literals appear in the **i-UIP** clause C_i , and increase their variable activity uniformly according to the current branching heuristic if the literals' variable activity have not yet been bumped during 1-UIP. The scheme does not bump variable activities for transient literals that are resolved away at non-conflicting level because, unlike transient literals at conflicting level, these literals cannot be re-asserted after the immediate backtracking. Notice that other post-analyze extensions such as Reason Side Rate (RSR) and Locality[] can be applied after applying **i-UIP-Inclusive** on C_i so that no literal's variable activity is double bumped.

$$\forall l \in C_i \cdot \text{NotBumped}(\text{var}(l)) \implies \text{bumpActivity}(\text{var}(l))$$

i-UIP-Exclusive collects literals appear exclusively in C_i and bump their variable activity uniformly. It also find all the literals in C_1 that are resolved away during **i-UIP**, and unbump their variable activity if they have been bumped during 1-UIP. The unbumped literals are no longer in the learned clause C_i , keep their variable activity bumped does not help solver to use C_i .

$$\begin{aligned} \forall l_i \in C_i \setminus C_1 \cdot \text{bumpActivity}(\text{var}(l_i)) \\ \forall l_1 \in C_1 \setminus C_i \cdot \text{unbumpActivity}(\text{var}(l_1)) \end{aligned}$$

3.5 Integrate with Chronological Backtracking

In SAT solver with chronological backtracking, literals on the assertion trail ϕ are not always sorted by decision levels. This change imposes a challenge to **i-UIP** learning since the previous implementation relies on solver’s ability to efficiently access all literals from any decision level in descending trail order (alg. 1 line 7).

To mitigate this challenge, we modify the solver to track the precise trail position of all asserted literals with a single vector. We then build a priority queue *lit_Order* to manage literal’s resolution order. The queue *lit_Order* prioritizes literals with higher decision level, and it favors literal with higher trail position when decision levels are tied. The order of *lit_Order* represents the correct resolution order of **i-UIP** because 1) an asserted literal’s decision level is the maximum level among all of its reasoning literals, and 2) an asserted literal always appears higher in the trail than its reasoning literal.

Alg. 4 is the pseudo-code implementation of the augmented **i-UIP** for chronological backtracking with the priority queue *lit_Order*. The algorithm first populates *lit_Order* with all literals in C_1 (line 4), and then continuously pops literals until the queue is empty (line 6). When a literal is resolved away, all of its reasoning literals are added into *lit_Order* if they are not already in the queue (line 13). The algorithm will always terminate because a literal cannot entered the queue twice (guaranteed by the properties of the trail order) and there are finite amount of literals on the trail.

4 Implementation and Experiments

4.1 Clause Reduction with i-UIP

To evaluate **i-UIP**’s effectiveness as a clause reduction technique, we implement **Pure-i-UIP** and **Min-i-UIP** on top of *MapleCOMSPS_LRB* [], the winner of SAT Race 2015 application track. We than compare the performance of the baseline *MapleCOMSPS_LRB* with Maple-**Pure-i-UIP** and Maple-**Min-i-UIP** on the full set of benchmarks from SAT RACE 2019 main track.

The benchmark contains 400 instances divided into two groups of 200, new and old, representing historical instances and fresh instances in the 2019 race,

Algorithm 4 i-UIP-CB**Require:** C_1 is a valid 1-UIP clause**Require:** ϕ is a valid assertion trail**Require:** lit_Order is a priority queue

```

1: procedure i-UIP-CB( $C_1, \phi, lit\_Order$ )
2:   ...
3:    $C_i \leftarrow C_1$  ▷ initialize i-UIP clause
4:   forall  $l \in C_1$  · Enqueue( $lit\_Order, l$ )
5:   DecisionLvs  $\leftarrow \{ level(l) \mid l \in C_1 \}$ 
6:   while  $lit\_Order \neq \emptyset$  do
7:      $p \leftarrow$  dequeue( $lit\_Order, l$ )
8:     if  $\exists q \in \text{Reason}(\neg p, \phi) \cdot level(q) \notin \text{DecisionLvs}$ 
9:        $\forall p$  is the last remanaing lit in its decision level then
10:        Pass
11:     else
12:        $C_i \leftarrow (C_1) \bowtie \text{Reason}(q, \phi)$ 
13:       forall  $l \in \text{Reason}(q, \phi) \cdot$  Enqueue( $lit\_Order, l$ )
14:     end if
15:   end while
16:   ...
17: end procedure

```

respectively. I partition the old group instances into six partitions of size 30 and one partition of size 20. Each partition is then assigned to a Intel CPU node with 16 cores (2 sockets 8 cores and 1 thread) and 96649 MB memory. The new group is partitioned based on their contributor (e.g. Heule contributed 22 matrix multiplication instances), and each partition is assigned to a aforementioned CPU node. To speed up the experiment, we allow a CPU node to solver at most seven instances concurrently.

Beside solved instances count and PAR-2 score, we additionally measure the average clause length and clause reduction ratio for each instances. For Maple-**Pure-i-UIP** and Maple-**Min-i-UIP**, we also captures the **i-UIP** learning attempted rate and success rate.

Solver	# solved	PAR-2	Clause Size	Cl Reduction%
<i>MapleCOMSPS_LRB</i>	221	5018.89	62.6	36.53%
Maple- Pure-i-UIP	226	4920.04	49.6	41.6%
Maple- Min-i-UIP	226	4890.67	45.2	47.8%

Fig. 1: Benchmark results of *MapleCOMSPS_LRB* , Maple-**Pure-i-UIP** and Maple-**Min-i-UIP** on SAT2019 race main track.

Fig. 1 shows that both version of **i-UIP** solved five more instances than the baseline solver with lower PAR-2 scores. **Min-i-UIP** has marginally lower

PAR-2 score than **Pure-i-UIP**. Both **Pure-i-UIP** and **Min-i-UIP** produce clause with significantly smaller size than 1-UIP by 27.7% and 20.7%, respectively. Fig. 2 shows the probability density distribution (PDF) of the average clause length from **Min-i-UIP** relative to 1-UIP. **Min-i-UIP** learning produces shorter clauses for 88.25% instances, and average relative reduction from 1-UIP is 18.685%. Fig. 3 compares the absolute average clause size from **Min-i-UIP** and 1-UIP, and it shows that **Min-i-UIP** in general produces smaller clauses, and the size reduction is more significant for instances with large average 1-UIP clause size.

We also looked at the 14 instances solved by **Min-i-UIP** but not by 1-UIP. **Min-i-UIP** produces smaller clauses for all of them with average relative reduction of 22% and maximum 77% (30 vs 135). Seven out of 14 instances has size relative reduction over 30%. For the 9 instances solved by 1-UIP but not by **Min-i-UIP**, **Min-i-UIP** only produce smaller clause by 33% and with average relative reduction of 3.3%.

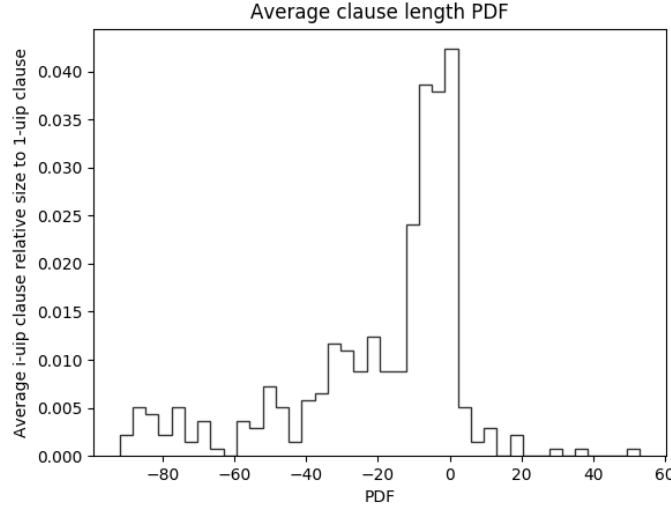


Fig. 2: Average clause (relative to 1-UIP clauses) size distribution. X axis indicates the relative size difference, and Y axis indicates the PDF.

Min-i-UIP outperformed **Pure-i-UIP** in both PAR-2 score and clause size. This results agrees with our observation in Fig. 4: **Min-i-UIP** attempted **i-UIP** learning more frequently, and it is more likely to succeed. Remark that the success of **i-UIP** learning is determined by the size of the learned i-UIP clause C_i , and the **i-UIP** learning frequency is also indirectly controlled by **i-UIP**'s success rate from the previous restart interval. The results indicates **Min-i-UIP**

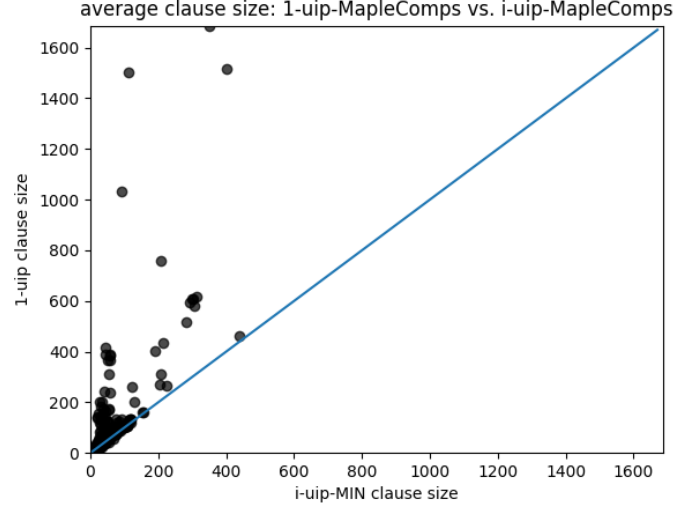


Fig. 3: Average clause size comparison plot. Each point in the plot represents an benchmark instance. X and Y axis shows the clause length from Maple-**Min-i-UIP** and *MapleCOMSPS_LRB*, respectively

shortened C_i 's size through further minimization of C_i at non-unique implication decision level.

Solver	i-UIP attempt rate	i-UIP success rate
Maple- Pure-i-UIP	16.1%	43.4%
Maple- Min-i-UIP	28.8%	59.3%

Fig. 4: Compare **Pure-i-UIP** and **Min-i-UIP** i-uip attempt rate and success rate. **Min-i-UIP** scheme attempted **i-UIP** more frequently, and it is more likely to successfully produce smaller C_i clause .

A solver produce smaller clauses can construct smaller proofs. For UNSAT instances, we additionally measure their DRUP [1] proof checking time as well as the size of the optimized DRUP proof. We used DART-trim [2] with 5000 timeout to check and optimize DRUP proofs.

Fig. 5 shows that the optimized proof construct by **Min-i-UIP** and **Pure-i-UIP** are significantly smaller than 1-UIP proofs. The relative proof size reduction roughly correlates to the average clause size reduction. Fig. 6 shows the absolute proof size comparison results.

Solver	optimized proof size (MB)	relative reduction size
<i>MapleCOMSPS_LRB</i>	613.9	0
Maple- Pure-i-UIP	487.2	6.90%
Maple- Min-i-UIP	413.2	17.18%

Fig. 5: Optimized UNSAT proof comparison for 1-UIP **Pure-i-UIP** and **Min-i-UIP**. Optimized proof size measures the average absolute proof size in MB, and relative reduction size measures the average relative reduction for all UNSAT instances.




Fig. 6: Average optimized proof size between 1-UIP and **Min-i-UIP**.

4.2 i-UIP as a Practical Learning Scheme

To evaluate **i-UIP**'s effectiveness as a clause learning scheme, we implement **Min-i-UIP** on *MapleCOMSPS_LRB* with the extensions mentioned in section ?? . We evaluated four different configurations (i-UIP, i-UIP-Greedy, i-UIP-Inclusive, and i-UIP-Exclusive) of **i-UIP** and 1-UIP learning on the SAT Race 2019 main track benchmark and report each configuration's solved instances, PAR-2 score and average clause size.

Fig. 7 summarizes the result of the experiment. Learning scheme i-UIP-greedy solves the same amount of instance (226) as i-UIP with less PAR-2 score. The inclusive activity adjustment solves the most SAT instances (138) and the least UNSAT instances (87). The exclusive activity adjustment scheme produces the shortest average clause size, but solved the second least instances, one more instance than the baseline.

Solver	# solved (SAT, UNSAT)	PAR-2	Avg clause Size
1-UIP	221 (132, 89)	5018.89	62.6
i-UIP	226 (135, 91)	4890.67	45.2
i-UIP-greedy	226 (135, 91)	4866.94	47.7
i-UIP-active-Inclusive	225 (138 , 87)	4958.49	52..12
i-UIP-active-Exclusive	223 (134, 89)	5015.23	43.2

Fig. 7: Benchmark results of 1-UIP (*MapleCOMSPS_LRB*), i-UIP(**Min-i-UIP**), i-UIP-Greedy, i-UIP-Inclusive, and i-UIP-Exclusive on SAT2019 race main track.

Inserted a table here, and graphs and analysis

4.3 i-UIP on Modern SAT solvers

To validate **i-UIP** as a generalizable learning scheme on modern SAT solvers, we re-implement **i-UIP** on *MapleLCMDist* [], *MapleLCMDiscChronoBT-DL-v3* [] and *expMaple_CM_GCBumpOnlyLRB*. The first two solvers are the winner of 2017 and 2019 SAT race, respectively. *expMaple_CM_GCBumpOnlyLRB* is a top ten solver from 2019 SAT race which uses random walk simulation to help branching. For each solver, we compare the base 1-UIP learning scheme against top two **i-UIP** configurations, **i-UIP-Greedy** and **i-UIP-Inclusive**, on the SAT Race 2019 main track benchmark. We report solved instances and PAR-2 score.

Inserted a table here, and graphs and analysis

References

1. LNCS Homepage, <http://www.springer.com/lncs>. Last accessed 4 Oct 2017